

SSID: Guest

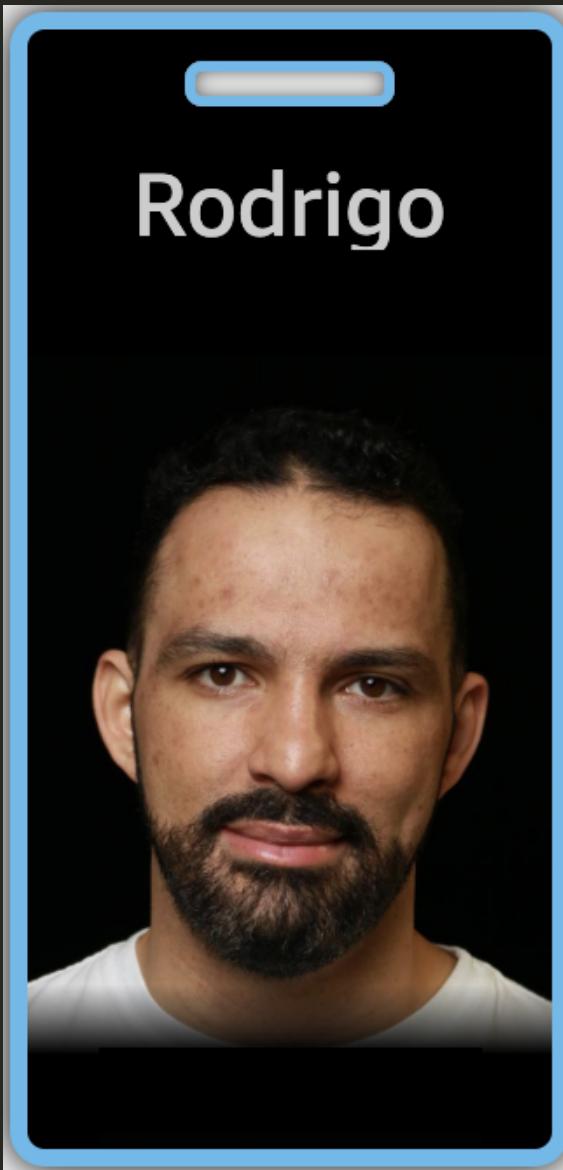
Password: BrokenWires@@2019

Getting Started with Kubernetes on AWS

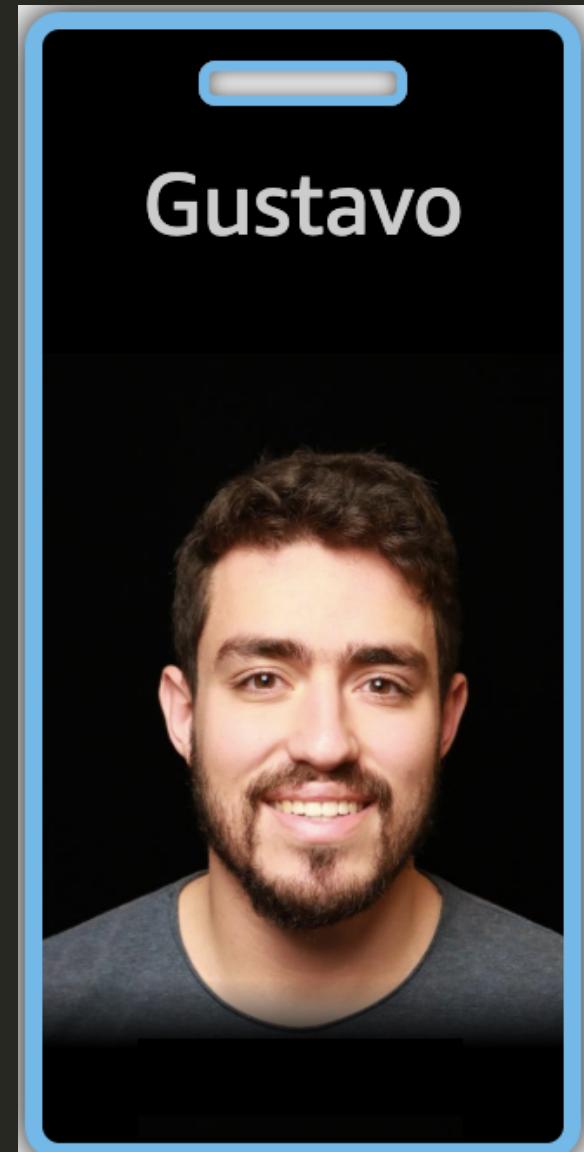
Brought to you by the AWS Cloud Support Team

Welcome and thank you!

Instructors



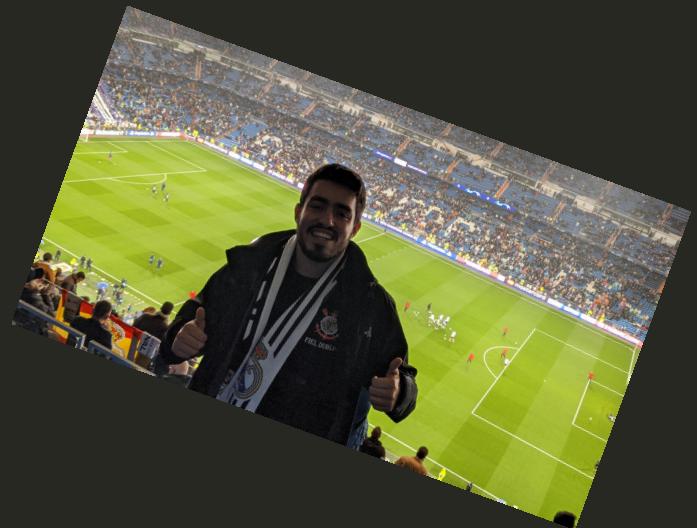
Rodrigo



Gustavo

About us

- Brazilian
- Football addicted and Corinthians supporter
- Love Italian Food :D



About us

- Brazilian from Brasília, the capital and one of the most beautiful cities in the world.
- A sports fan, former handball player, and Corinthians supporter.
- A language enthusiast. I speak English, Portuguese, Spanish and a bit of French. It would be great if you guys could teach me some Italian during the intervals.

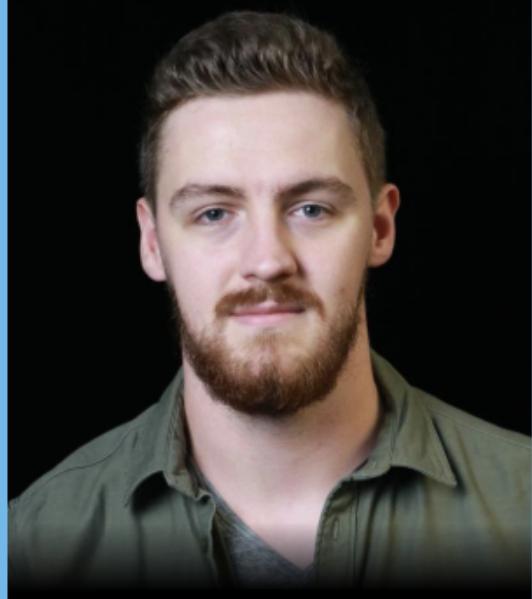


Helpers around the room

Jonathan



Sebastian



What is Cloud Support Engineering?

- Cloud Support Engineers are the front line of AWS' technical support
- Multiple teams, each with specific focus areas.
- What do we do? A surprising amount:
 - Customer interactions
 - Training
 - Learning is part of the job
 - Hiring

What do we do as Cloud Support Engineers - Deployment

- We assist customers to develop services and technologies built on top of AWS Cloud Platform.
 - ECS and EKS (container orchestration)
 - CloudFormation, Elastic Beanstalk, OpsWorks (Automating operations with Chef)
 - CodeDeploy, CodePipeline, CodeCommit (Implementing CI/CD pipelines on AWS)
- Apart from working on a broad spectrum of technical issues, we also work actively:
 - Coaching and mentoring new hires.
 - Developing and delivering trainings
 - Recruiting, interviewing and participating on the hiring process.

The Team



Day 1

Agenda Day1

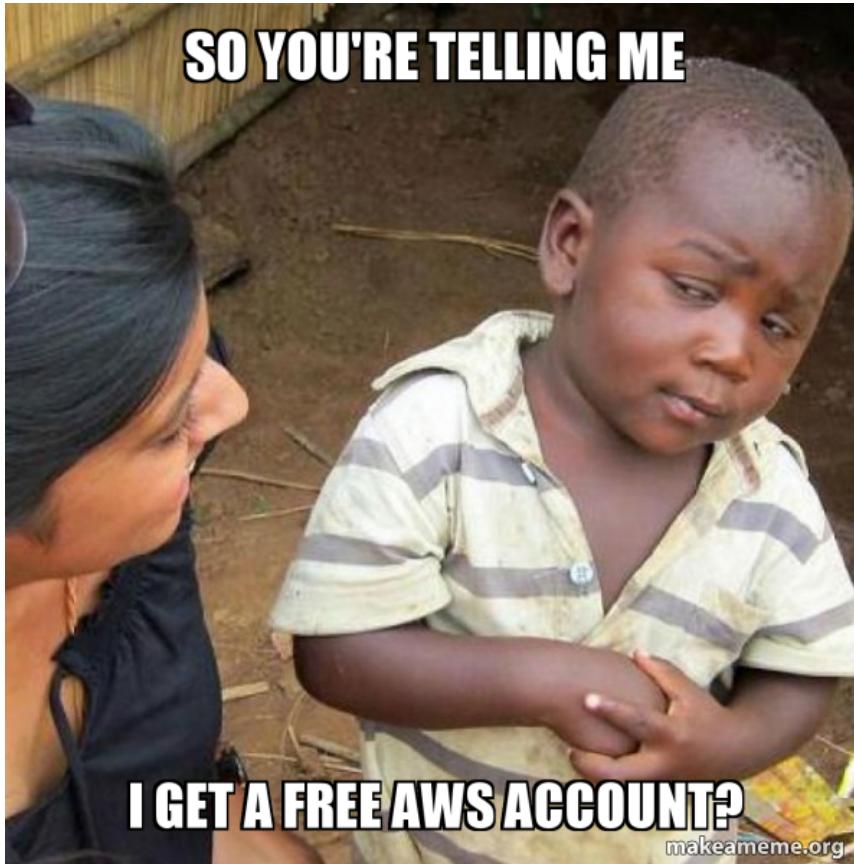
- Introduction to Containers
- Docker Overview
- Kubernetes Explained
- EKS - Exploring and Deploying
- Environment Setup
 - Break (Networking and Food)
- Exploring the cluster with kubectl
- Deploying an application with Kubernetes
- Kubernetes Services

Agenda Day2

- What have we learnt?
- Review questions.
- Challenge

Before we get started

Setting up AWS Account and Environment



Login into your AWS Account

1. Navigate to: <https://dashboard.eventengine.run/login>
2. Enter the team hash provided to you on arrival.
3. Click the AWS Console button.
4. Click the Open AWS Console button to log in to the account.
5. Optionally, credentials are provided for CLI access.

Note

The account will be deactivated at the end of the Event.

Setting up the environment

Launching your Lab Environment

1. Cloud9 - this will be where you'll be performing the labs throughout the sessions.

Follow the steps provided in the README.md to setup your environment

<https://github.com/aws-els-lin/eks>



Containers and Docker

Container Overview

- What is a container?

Container Overview

- What is a container?
- What is Docker? what benefits does it provide?

Container Overview

- What is a container?
- What is Docker? what benefits does it provide?
 - Consistent environment

Container Overview

- What is a container?
- What is Docker? what benefits does it provide?
 - Consistent environment
 - Portability of code, runtime, system tools, system libraries, etc.

Container Overview

- What is a container?
- What is Docker? what benefits does it provide?
 - Consistent environment
 - Portability of code, runtime, system tools, system libraries, etc.
 - Version control

Container Overview

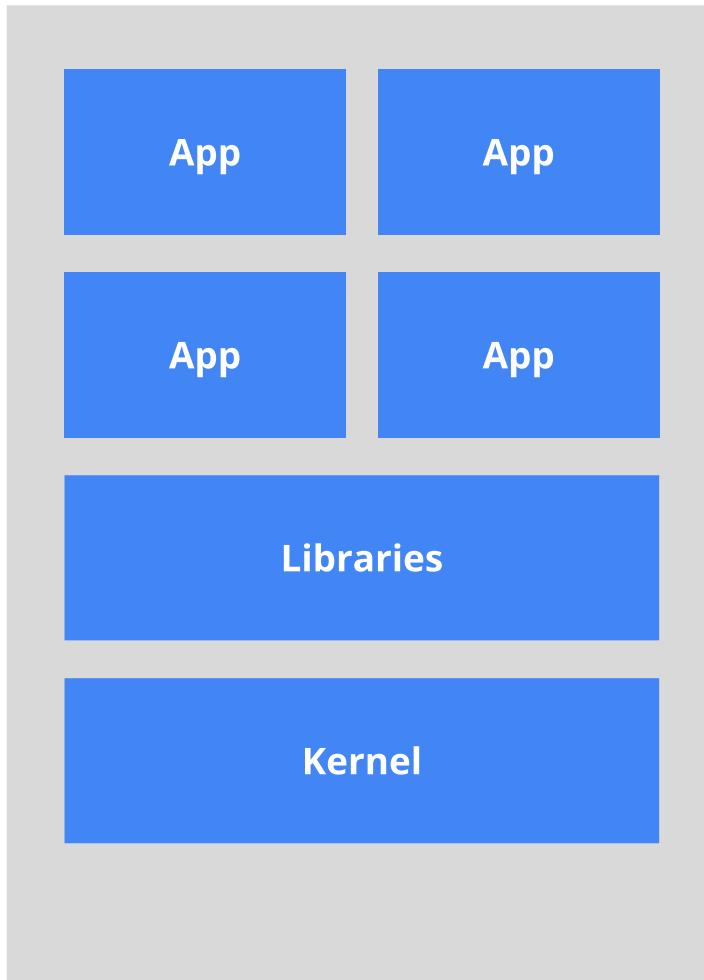
- What is a container?
- What is Docker? what benefits does it provide?
 - Consistent environment
 - Portability of code, runtime, system tools, system libraries, etc.
 - Version control
 - Lightweight

Container Overview

- What is a container?
- What is Docker? what benefits does it provide?
 - Consistent environment
 - Portability of code, runtime, system tools, system libraries, etc.
 - Version control
 - Lightweight
 - Microservices

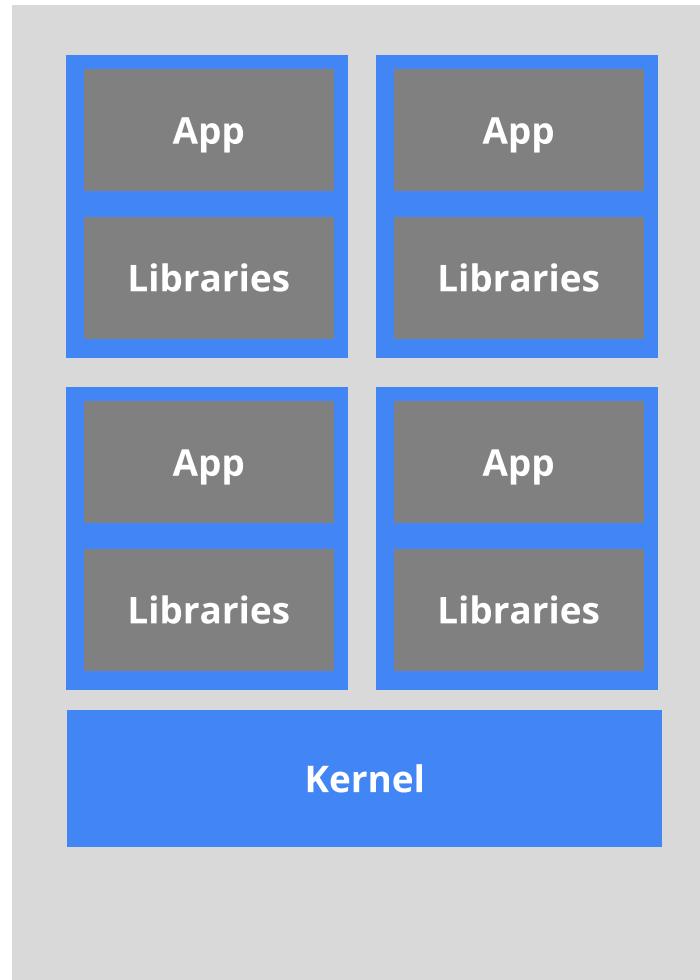
Why containers

The old way: Applications on host



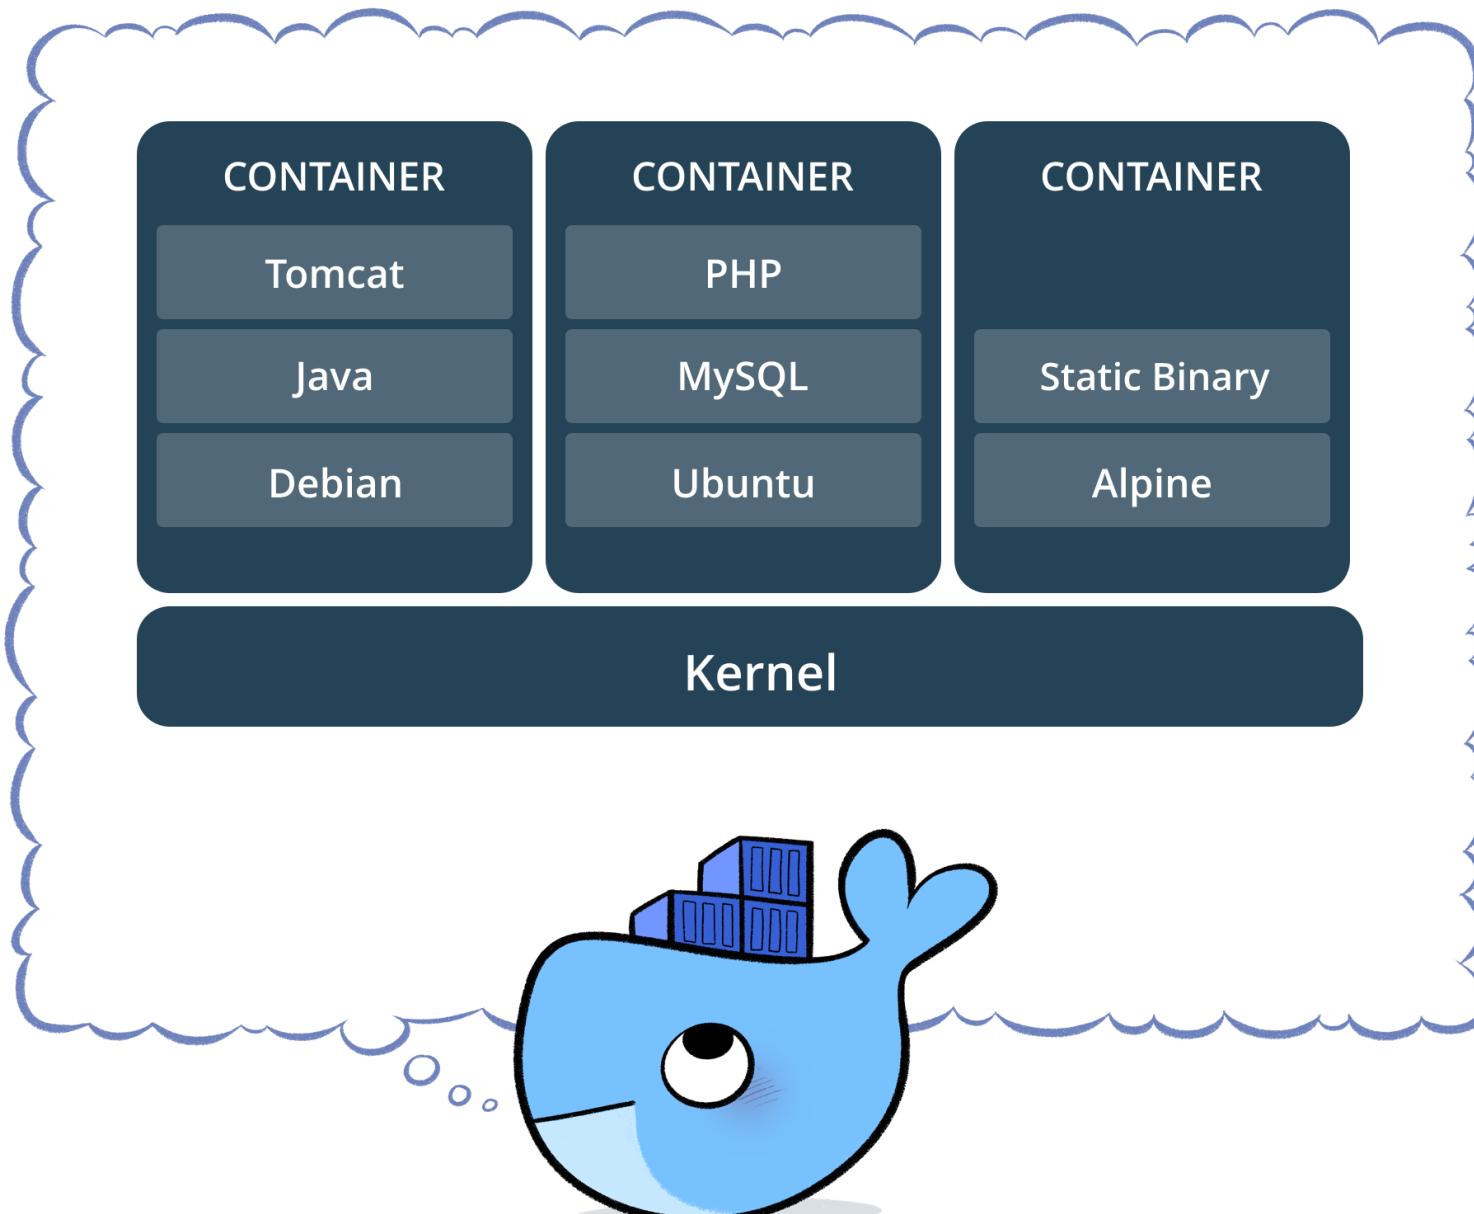
*Heavyweight, non-portable
Relies on OS package manager*

The new way: Deploy containers

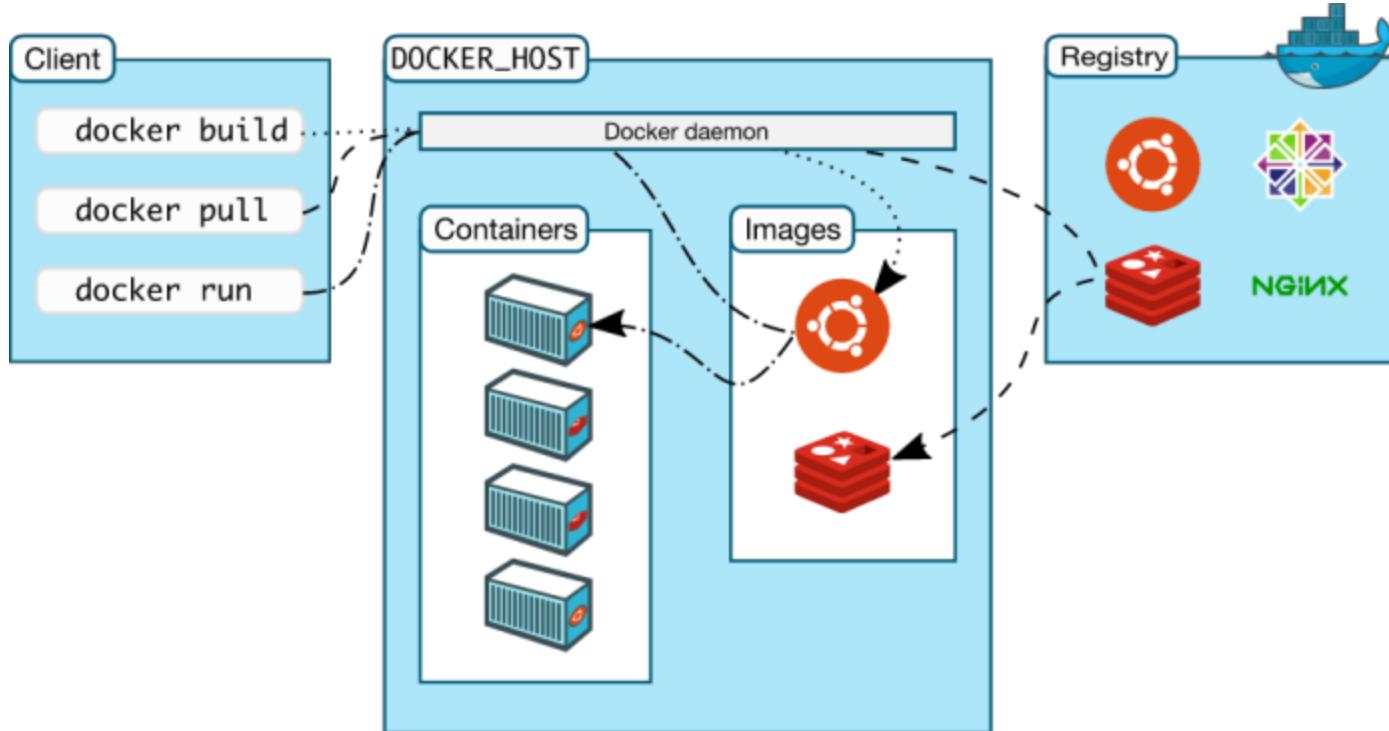


*Small and fast, portable
Uses OS-level virtualization*

Container Overview



Container Lifecycle



Docker commands - pull

Pull an image from an image registry

```
Admin:~/environment $ 
```

Docker commands - run

Running a container from an image

```
Admin:~/environment $ █
```

The `-p` flag maps a port from the host to the container

The `-d` flag runs the container in the background - detached

Let's docker

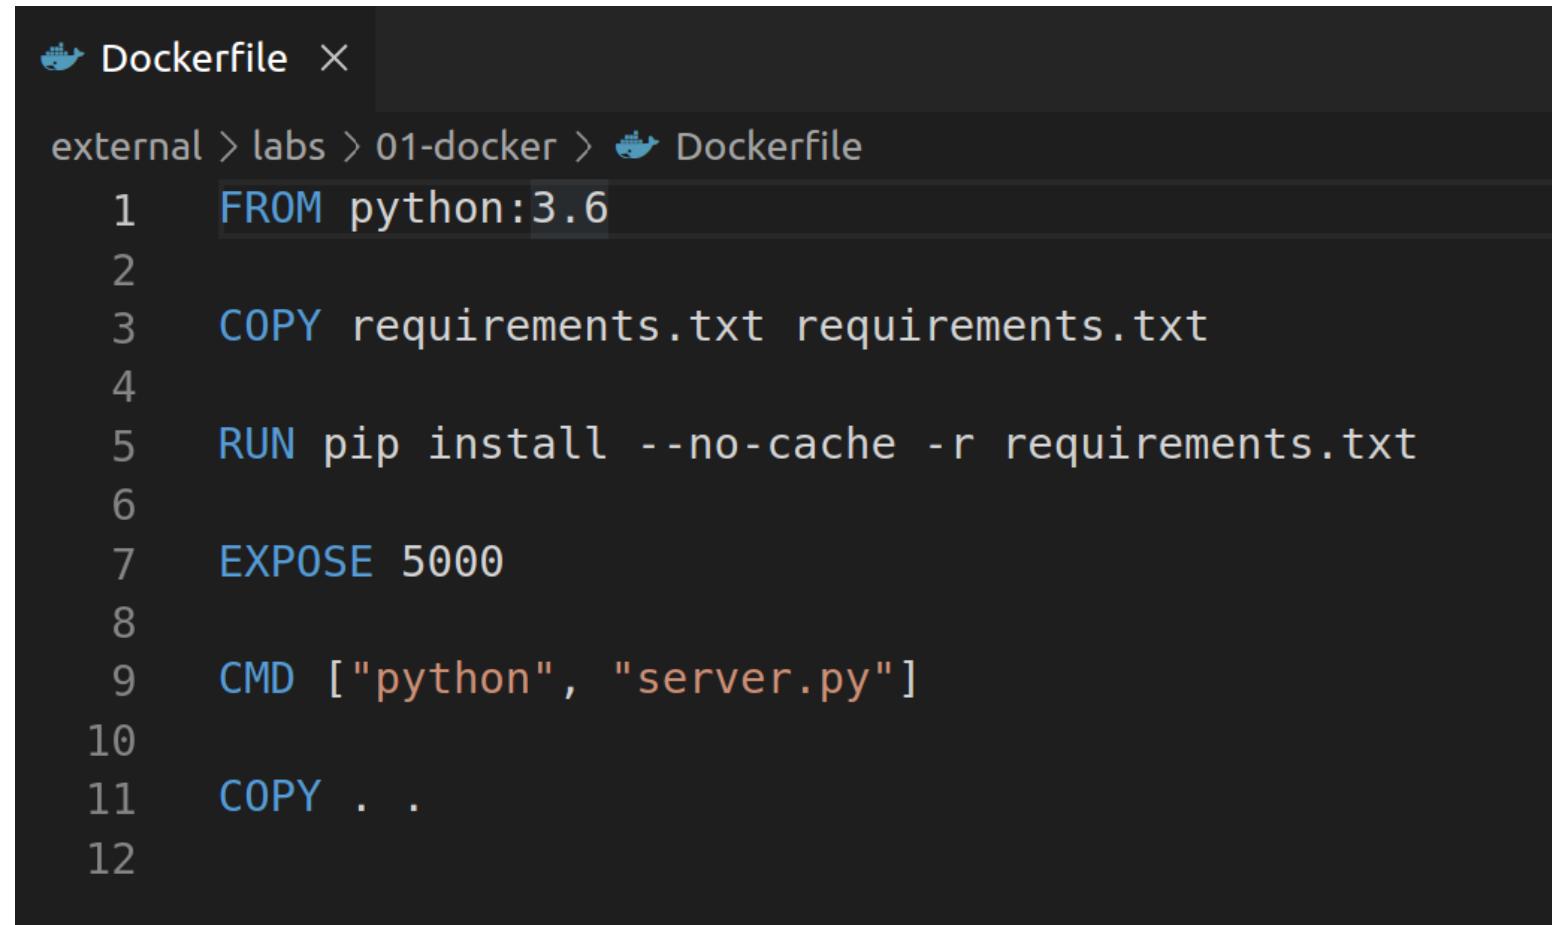
Lab 1: Building & Running a Docker Image

<https://github.com/aws-els-lin/eks/tree/master/labs/01-docker>

Dockerfile

Dockerfile contains a list of instructions to build a container image

```
labs/01-docker/Dockerfile
```



The screenshot shows a code editor window with a dark theme. The title bar says "Dockerfile X". The file path is "external > labs > 01-docker > Dockerfile". The code itself is a Dockerfile with the following content:

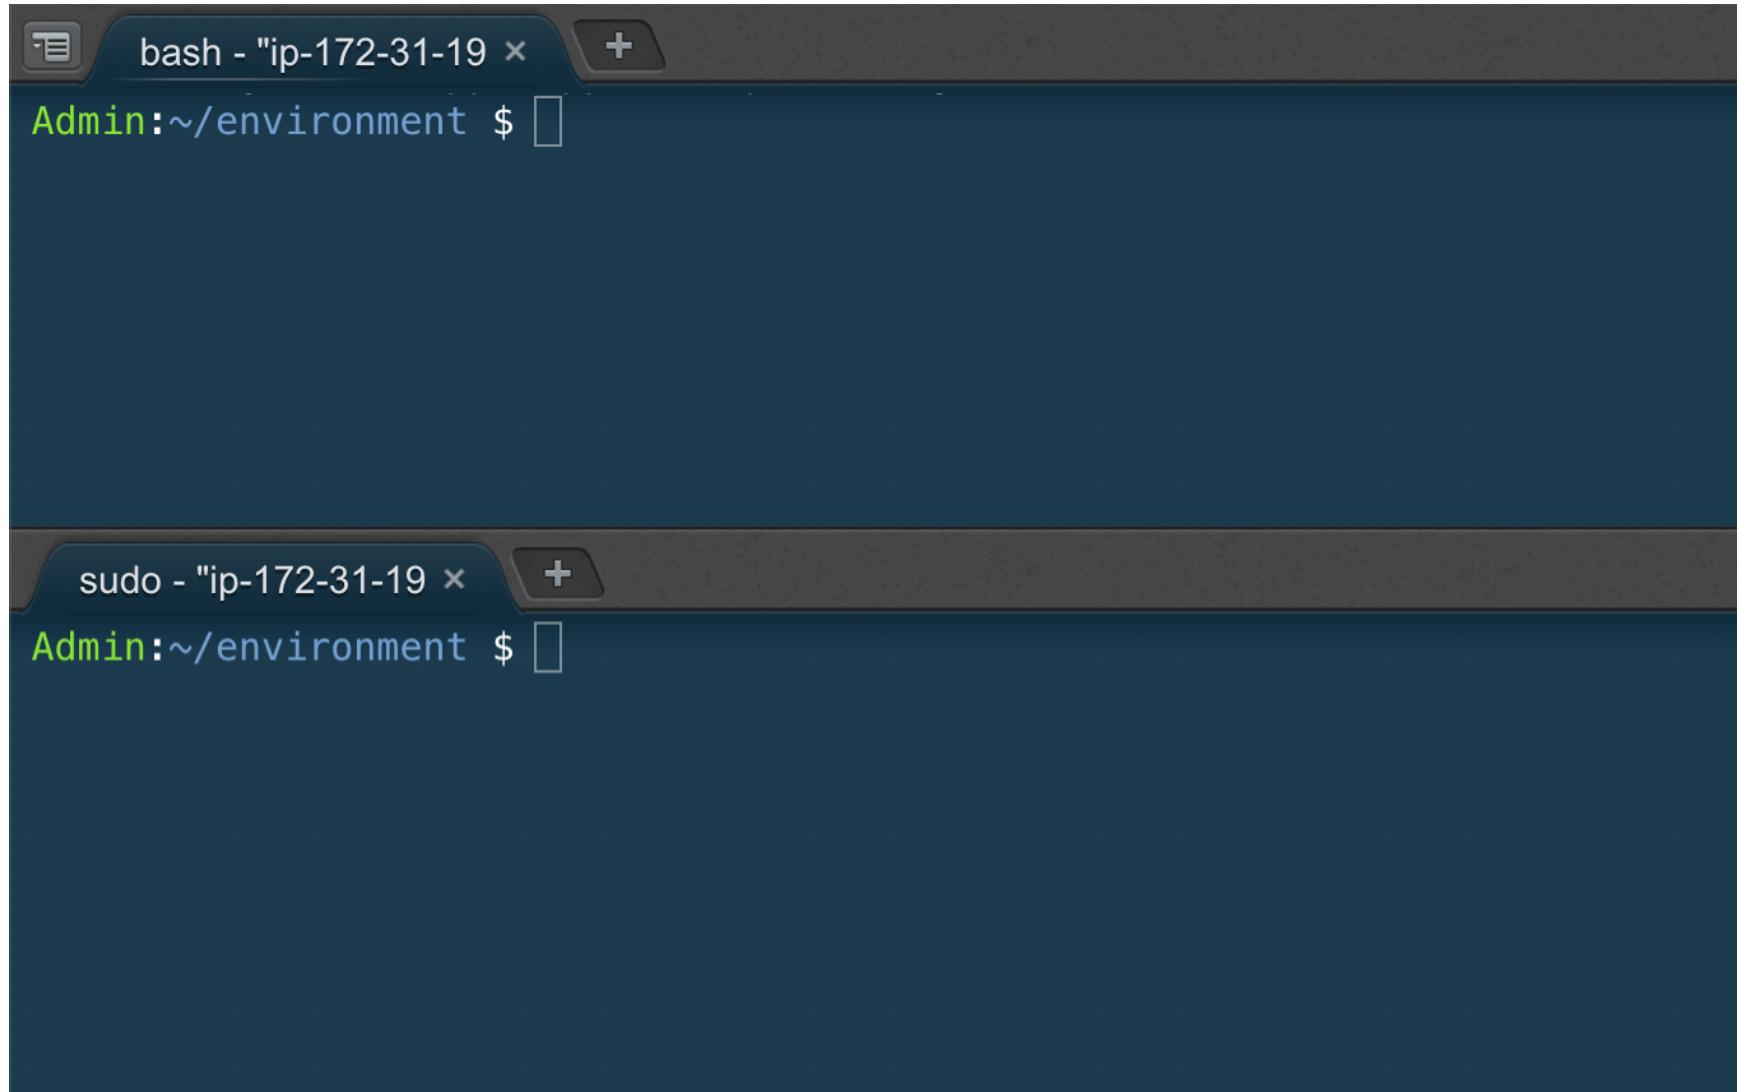
```
1 FROM python:3.6
2
3 COPY requirements.txt requirements.txt
4
5 RUN pip install --no-cache -r requirements.txt
6
7 EXPOSE 5000
8
9 CMD ["python", "server.py"]
10
11 COPY . .
12
```

Docker build

```
Admin:~/environment/python-app $ █
```

Docker run

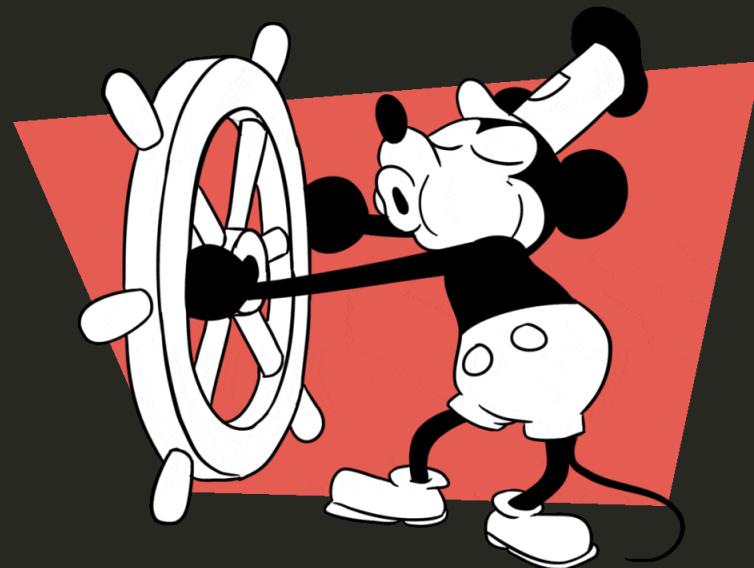
Finally we can run a container from the image



The image shows a terminal window with two tabs. The top tab is titled "bash - "ip-172-31-19 x"" and the bottom tab is titled "sudo - "ip-172-31-19 x"". Both tabs show the prompt "Admin:~/environment \$". The terminal has a dark blue background and a light gray header bar.

Before we move on, any questions?

What is



Kubernetes?

What is Kubernetes?



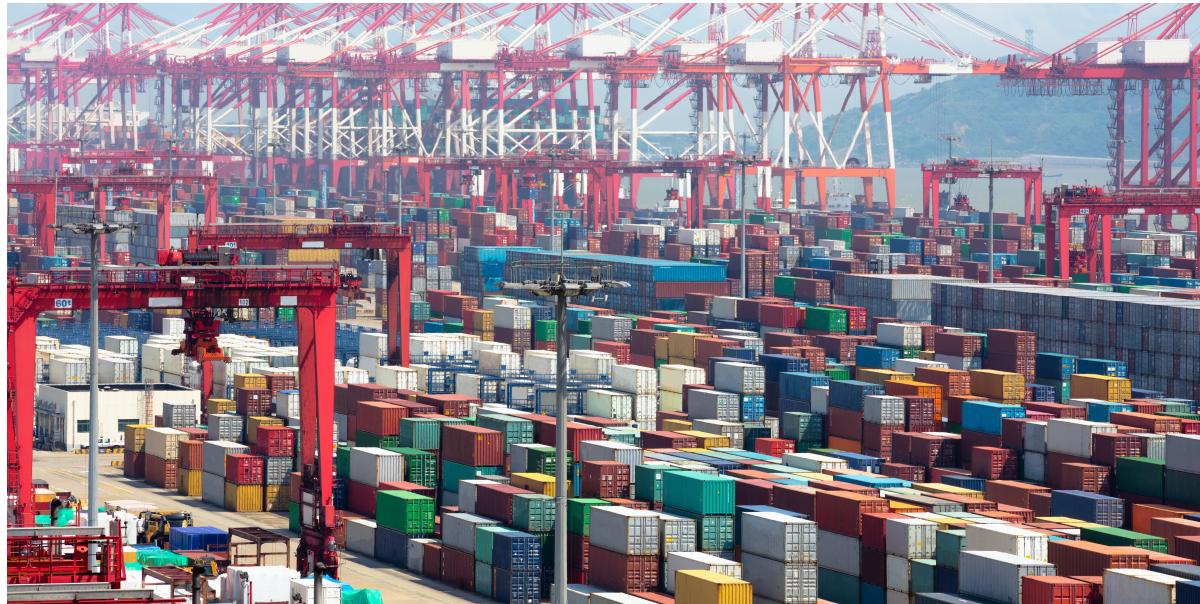
Why Kubernetes?

Why Kubernetes?



Why Kubernetes?

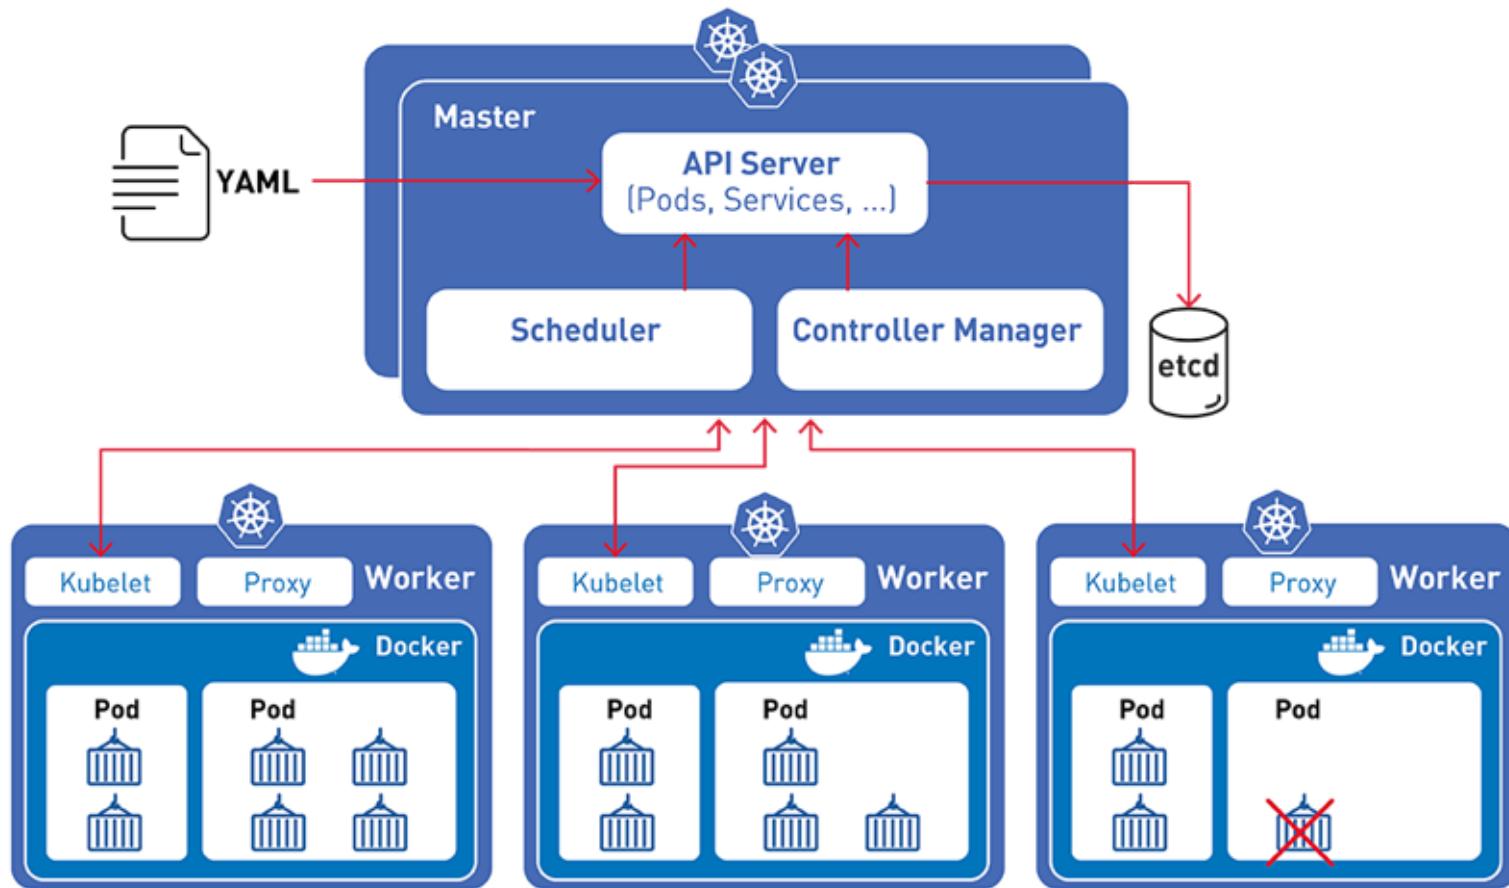
- Containers are a good way to bundle and run your applications.
- Containers paved the way to build cloud native systems, in which services are implemented using small clouds of containers.
- However, in a production environment, you need to manage the thousands of containers that run the applications and ensure that there is no downtime.



Key capabilities were missing !

- Resource Utilization
- Using multiple containers with shared resources
- Monitoring running containers
- Handling dead containers
- Autoscaling container instances to handle load
- Making the container services easily accessible Connecting containers to a variety of external data sources

Kubernetes Architecture





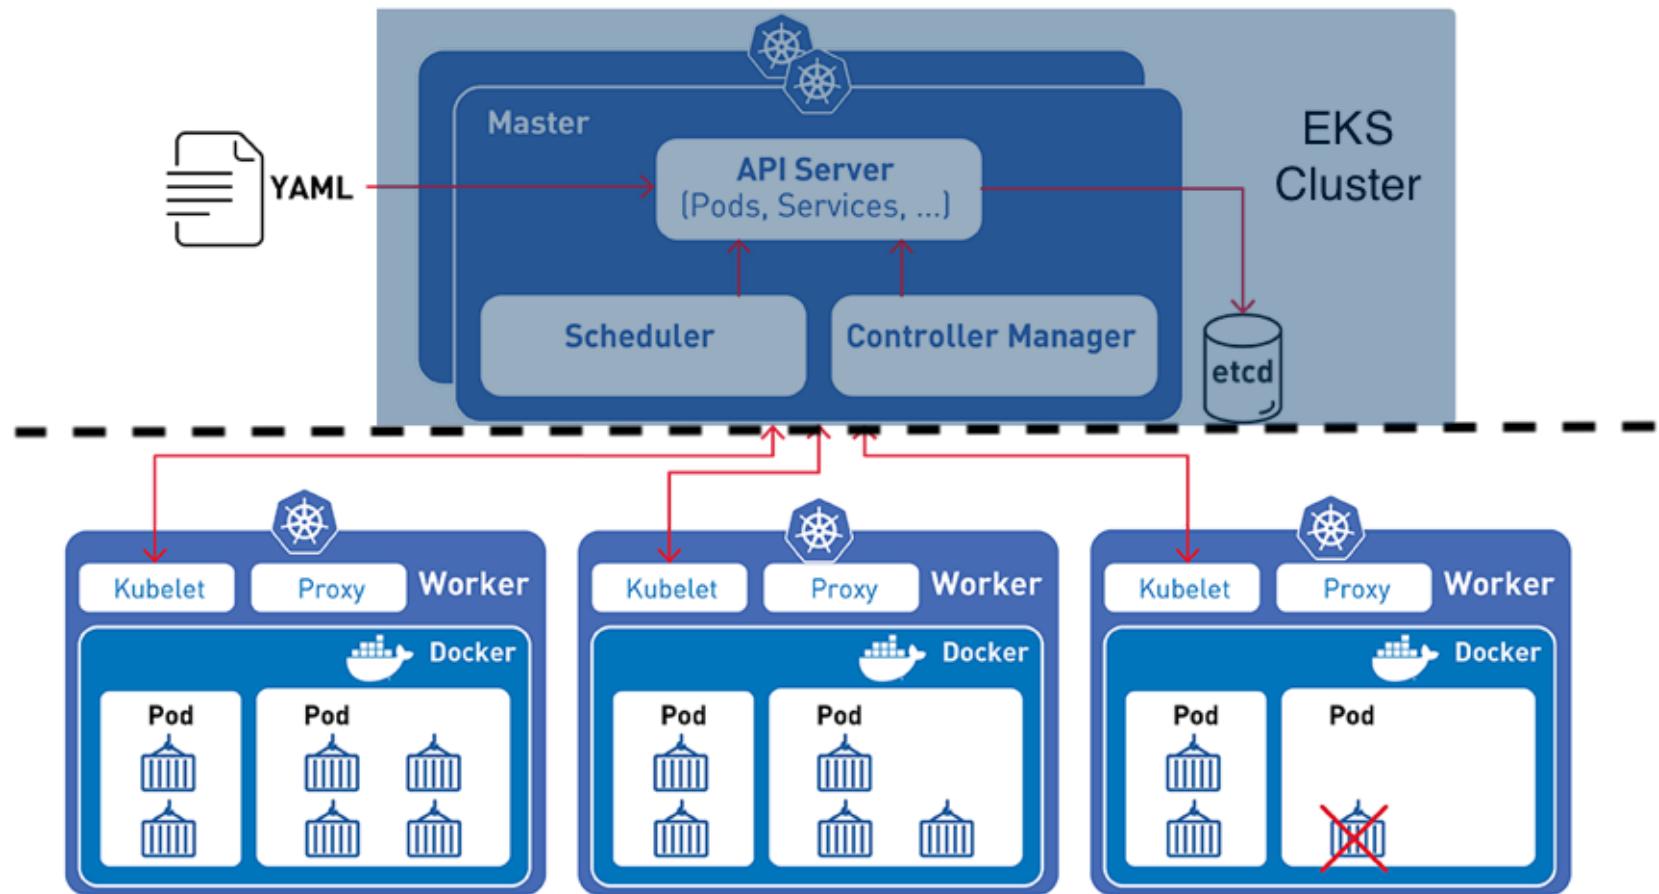
EKS

What is EKS?

What is EKS?

- EKS is the **managed** version of Kubernetes offered by AWS
- EKS launches and maintains the Control Plane for you with **high available** components
- EKS offers integration with other AWS services such as VPC and IAM. These integrations are **open source** projects built with the community
- EKS takes care of upgrades and patching
- EKS is based on vanilla Kubernetes

What is EKS?



Connecting to your K8s/EKS Cluster

What we'll need

- A client, which is a single binary called `kubectl`

What we'll need

- A client, which is a single binary called `kubectl`
- A configuration file for kubectl to be stored under: `~/.kube/config`

What we'll need

- A client, which is a single binary called `kubectl`
- A configuration file for kubectl to be stored under: `~/.kube/config`

The kubectl configuration can be generated automatically.

Launching a Kubernetes Cluster / EKS

Creating a Cluster

For this we'll use `eksctl` to create an EKS cluster and Worker Nodes:

```
eksctl create cluster --version 1.14 --node-type t3.medium --name eks
```

This command is also available on the [GitHub page](#)

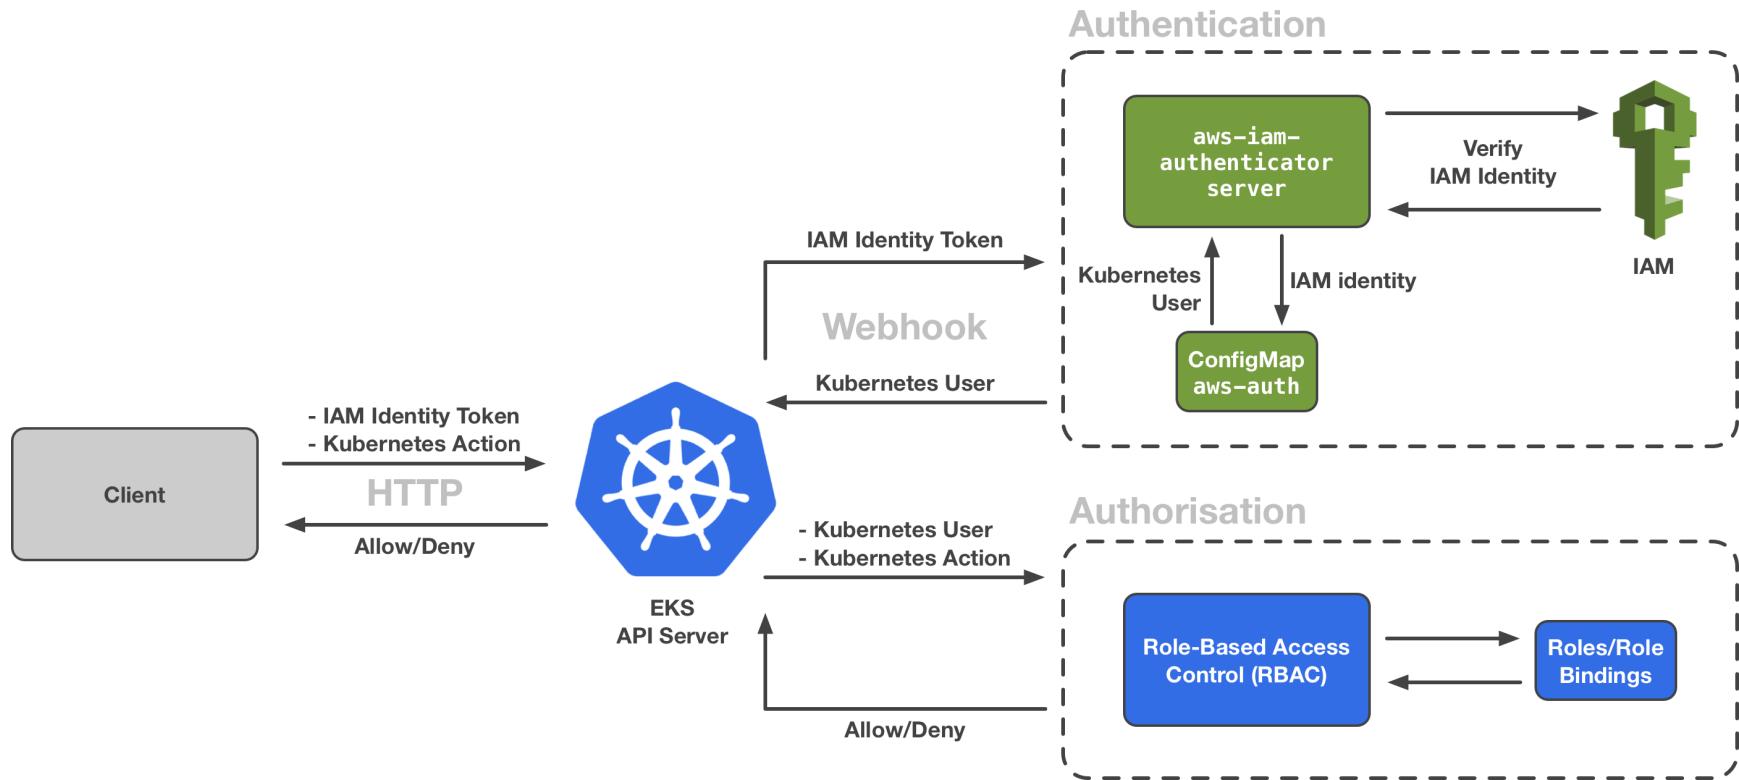
The EKS Cluster will take approximately 15 minutes to create

While that runs...

kubectl configuration - eksctl

```
Admin:~/environment $ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+0MITTED
    server: https://07413FD8A52FEC4D8C73360C03988541.yl4.ap-southeast-1.eks.amazonaws.com
    name: eks.ap-southeast-1.eksctl.io
contexts:
- context:
    cluster: eks.ap-southeast-1.eksctl.io
    user: i-0d8228356c7f829a6@eks.ap-southeast-1.eksctl.io
    name: i-0d8228356c7f829a6@eks.ap-southeast-1.eksctl.io
current-context: i-0d8228356c7f829a6@eks.ap-southeast-1.eksctl.io
kind: Config
preferences: {}
users:
- name: i-0d8228356c7f829a6@eks.ap-southeast-1.eksctl.io
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
      - token
      - -i
      - eks
      command: aws-iam-authenticator
      env: null
```

kubectl - Authentication



configmap - Authentication

Check to see if you have already applied the `aws-auth` ConfigMap.

```
kubectl describe configmap -n kube-system aws-auth
```

configmap - Authentication

Example ConfigMap:

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
data:
  mapRoles: |
    - rolearn: arn:aws:iam::555555555555:role/devel-worker-nodes-NodeInstanceRole-74RF4UBDUKL6
      username: system:node:{EC2PrivateDNSName}
      groups:
        - system:bootstrappers
        - system:nodes
  mapUsers: |
    - userarn: arn:aws:iam::555555555555:user/admin
      username: admin
      groups:
        - system:masters
    - userarn: arn:aws:iam::111122223333:user/ops-user
      username: ops-user
      groups:
        - system:masters
```

Let's check our new Cluster

Let's check our new Cluster

1 - Head to the EKS service in the AWS Web Console

Let's check our new Cluster

1 - Head to the EKS service in the AWS Web Console

2 - Select the 'eks' cluster and check what options are available

Checking your EKS cluster

The screenshot shows the AWS EKS Cluster configuration page for a cluster named 'eks'. The page is divided into sections: General configuration, Networking, and a bottom navigation bar.

General configuration:

- Kubernetes Version: 1.14
- Platform Version: eks.1
- Status: ACTIVE
- API server endpoint: https://CA124C4EA7486B52460C551DE1787C05.yl4.eu-north-1.eks.amazonaws.com (highlighted with a red box)
- OpenID Connect provider URL: https://oidc.eks.eu-north-1.amazonaws.com/id/CA124C4EA7486B52460C551DE1787C05
- Cluster ARN: arn:aws:eks:eu-north-1:335688126910:cluster/oscar
- Certificate authority: A long string of characters (highlighted with a yellow box): LS0tLS1CRUdjTIBDRVJUSUZJQ0FURS0tLS0tCk1JSUN5RENDQWJDZ0F3SUJBZ0lCQURBTkJna3Foa2lHOXcwQkFRc0ZBREFWTVJNd0VRWURWUVFERXdwcmRXSmwKY201bGRHvnPNojRYRFRNU1Ea3lNekv3TwPjMU1Wb1hEVek1TURreU1ERXdNakkxTVZvd0ZURVRNQKVHQTFVRQpBeE1LYTNWaVpYSnVaWFJsY3pDQOFtSYdElIV1KSC29aSWh2Vn5R1II1IV1C01FBRGd0pVBBRENDD0V1EvO2dnR1IIR50w0CmbhN1R2sTIII3MA01FERD
- Role ARN: arn:aws:iam::335688126910:role/eksctl-eks-cluster-ServiceRole-1X058FD0AGYIO

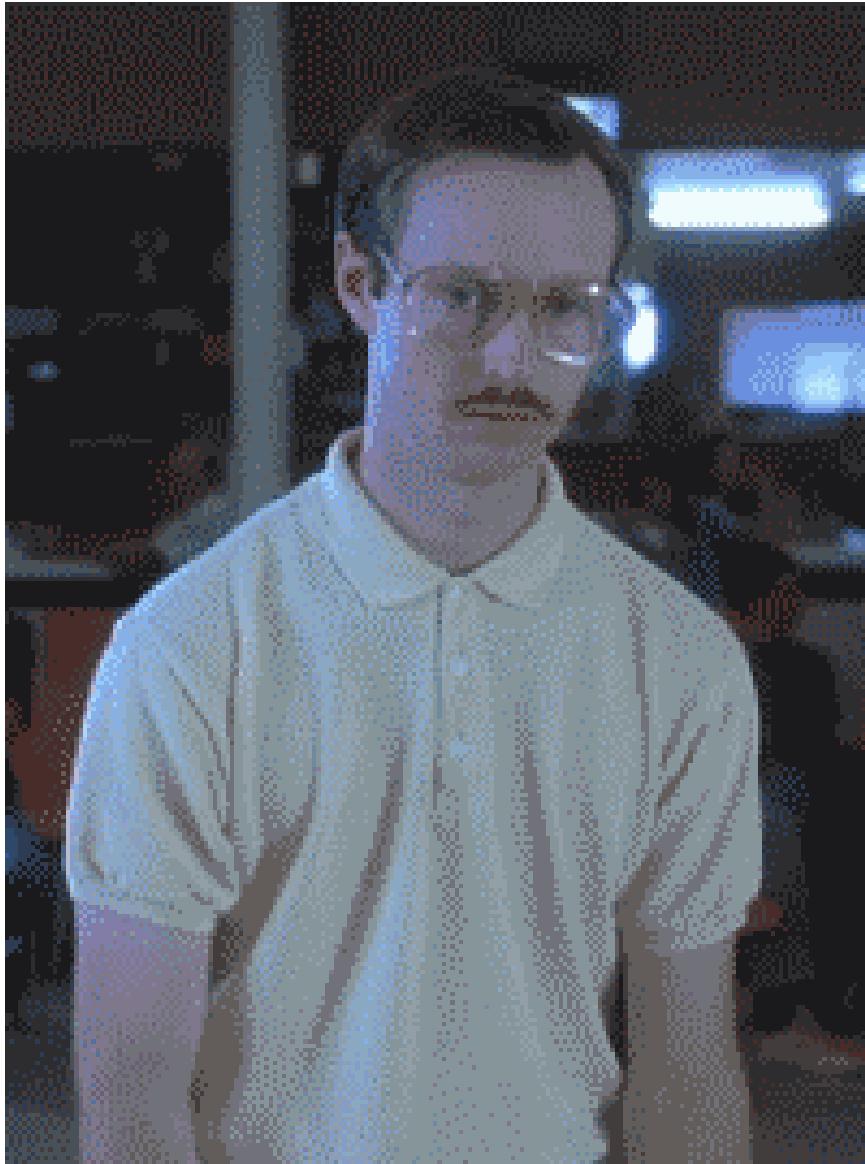
Networking:

VPC	Subnets	Security groups	API server endpoint access
vpc-04287be37e224d2bc	subnet-09a905b81dc255e44 subnet-0ffae02dd45bc1ed6 subnet-0a8504adb33a2a814 subnet-09d0feb73ac9f0f98 subnet-0d239c4720e4cd59f subnet-019a941b7efc42ce4	sg-0258eba705597213e	Private access Disabled

Bottom navigation:

- eks
- Cluster configuration
- Nodes
- Logs
- Events
- Metrics
- Logs Insights
- Logs Metrics
- Logs Metrics Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights Insights Insights Insights
- Logs Metrics Insights Insights Insights Insights
- Logs Metrics Insights Insights
- Logs Metrics Insights
- Logs Metrics
- Logs

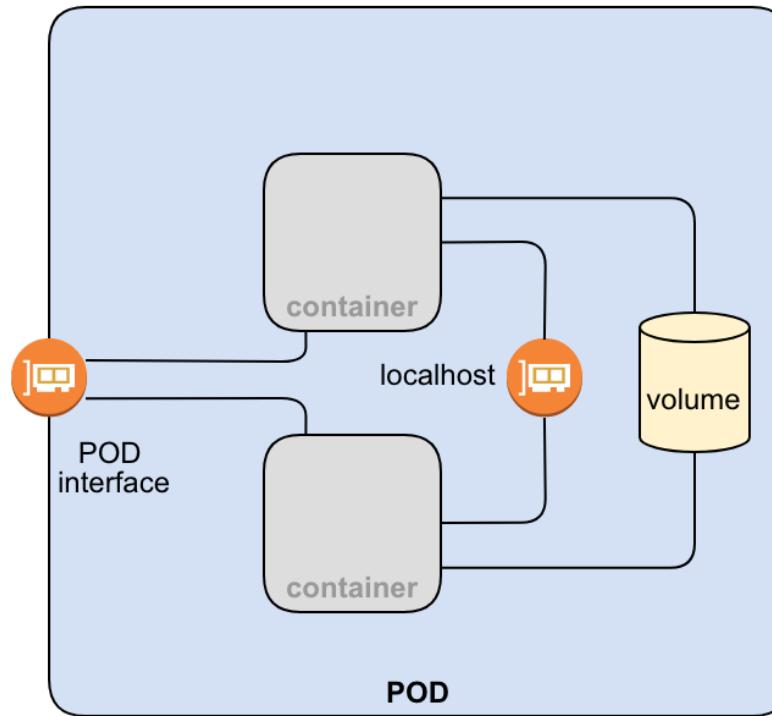
Successfully Create and Connect to EKS Cluster



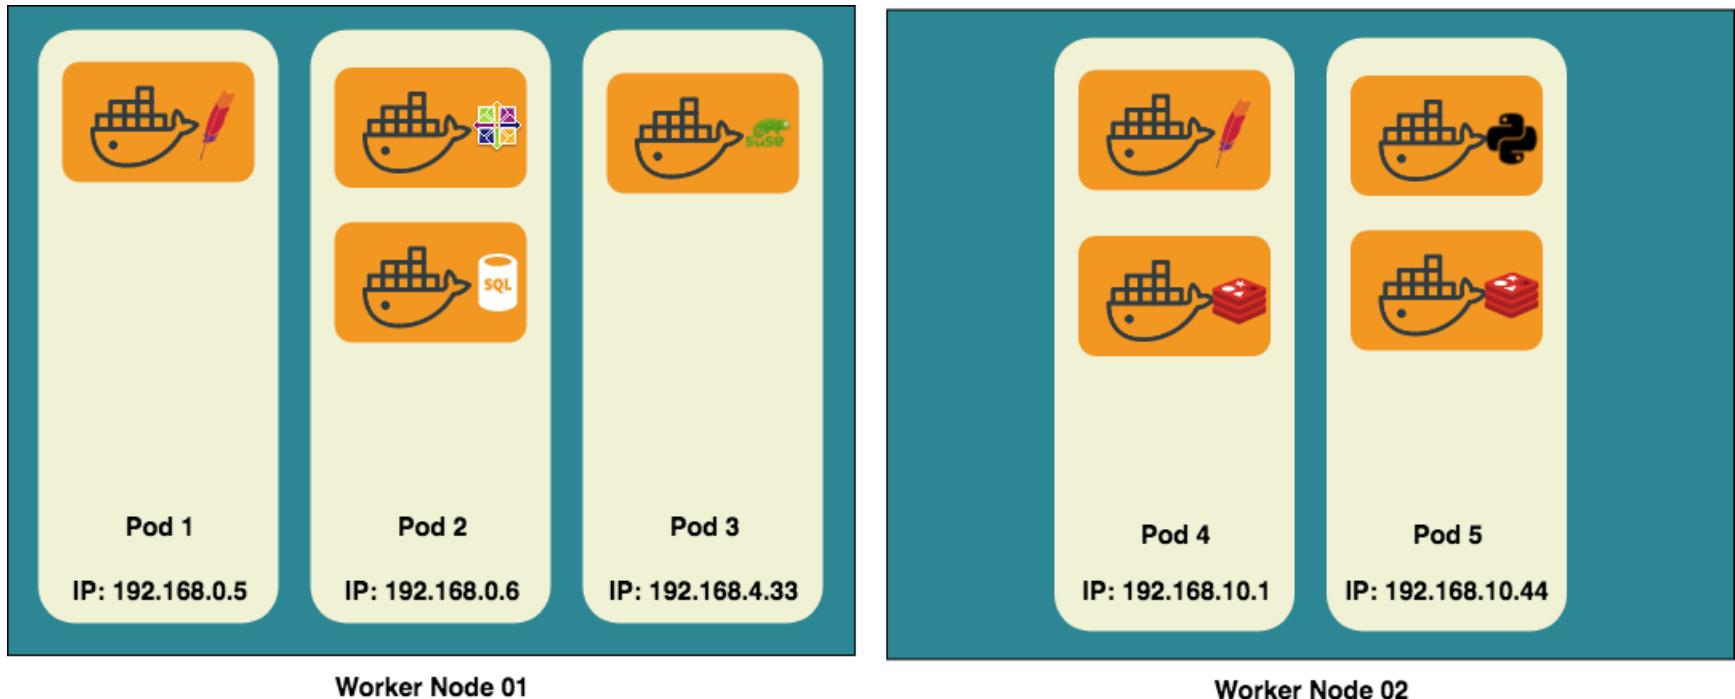
Pod

What is a Pod?

- The smallest building block of Kubernetes
- A Pod encapsulates the container(s) and resources needed to run the application
- A unit of deployment



What is a Pod?



Creating Pods in Kubernetes

Let's kubectl

Lab 2: Introduction to Pods

<https://github.com/aws-els-lin/eks/tree/master/labs/02-pods>

Lab 2: Define a Pod

Pod definition

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
    - name: container1
      image: nginx
```

Lab 2: Creating a Pod

Send the definition to the cluster:

```
$ kubectl apply -f pod.yaml  
pod/web-server created
```

Lab 2: Check the Pod

List and describe the Pod

```
# View the deployed pod
$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
web-server            1/1     Running   0          10s

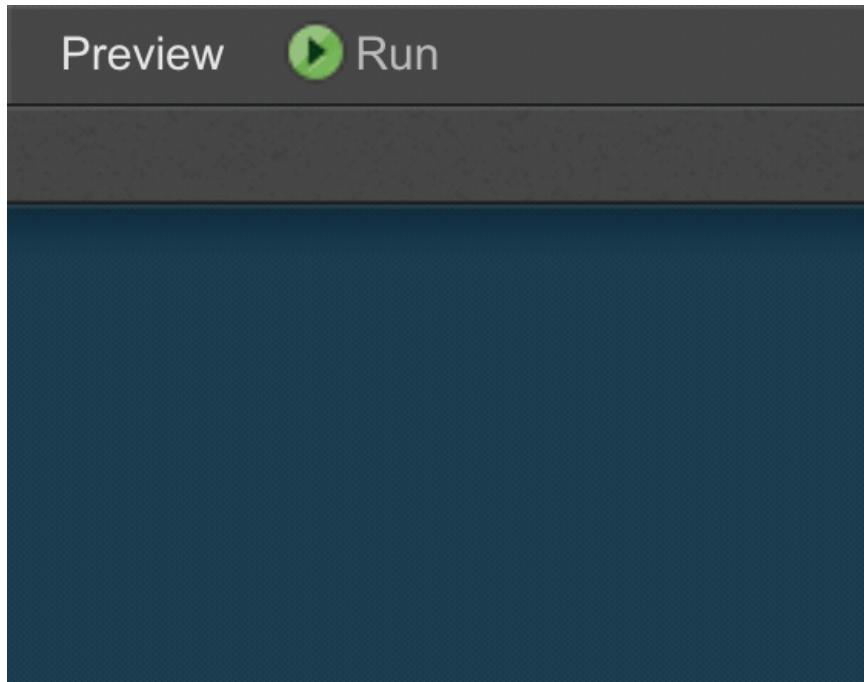
# See details of the pod
$ kubectl describe pod web-server
Name:                 web-server
Namespace:            default
Priority:             0
PriorityClassName:   <none>
Node:                 ip-10-0-100-30.eu-north-1.compute.internal/10.0.100.30
Start Time:           Fri, 27 Sep 2019 15:55:35 +0200
Labels:               <none>
[...]
```

Lab 2: Check the Pod

Create a tunnel and connect to your Pod

```
kubectl port-forward pod/web-server 8080:80 &  
curl localhost:8080
```

Or, connect using a browser and Cloud9



Lab 2: Clean up

We ran the port-forward in the background, lets clean it up before we move on

In the Cloud9 terminal to bring it back to the foreground:

```
fg
```

```
# Hit Ctrl + C to kill the port-forward
```

Working with Pods

Lab 3: Working with Pods

Check the logs

```
$ kubectl logs web-server
```

Lab 3: Working with Pods

Check the logs

```
$ kubectl logs web-server
```

Connect

```
# run one command  
kubectl exec web-server cat /etc/hostname  
  
# run with console connected to pod  
kubectl exec -it web-server -- bash
```

Lab 3: Working with Pods

Check the logs

```
$ kubectl logs web-server
```

Connect

```
# run one command  
kubectl exec web-server cat /etc/hostname  
  
# run with console connected to pod  
kubectl exec -it web-server -- bash
```

You can delete it: *But keep it for now!*

```
kubectl delete pod web-server  
  
# OR  
  
kubectl delete -f labs/01_pod.yaml
```

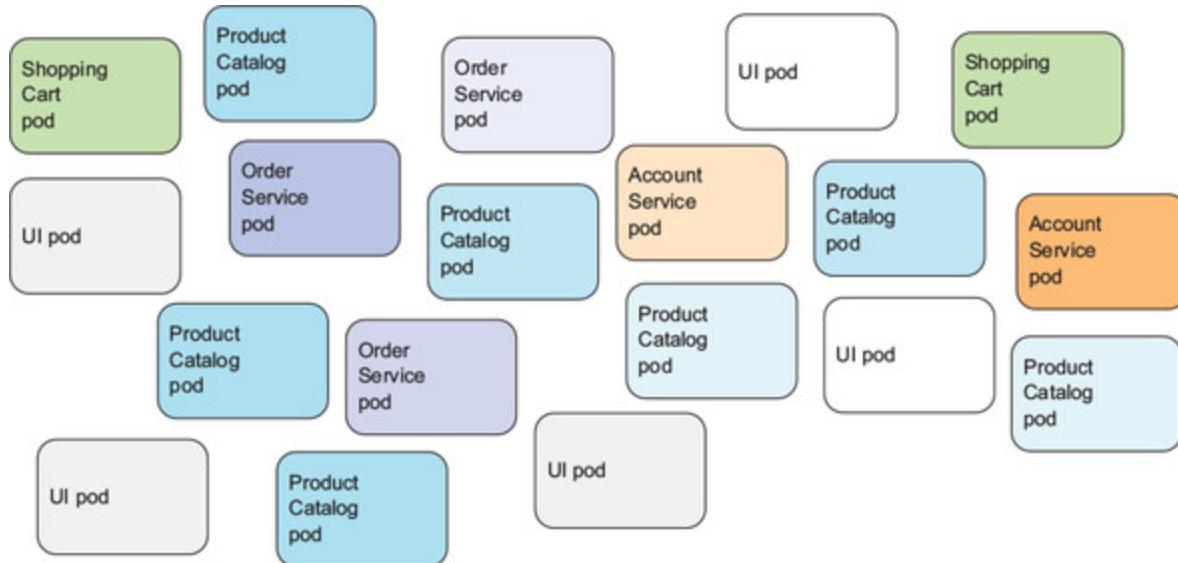
Let's kubectl

Lab 3: Playing around with our Pod

<https://github.com/aws-els-lin/eks/tree/master/labs/03-more-pods>

Labels

What if we are running a lot of pods?

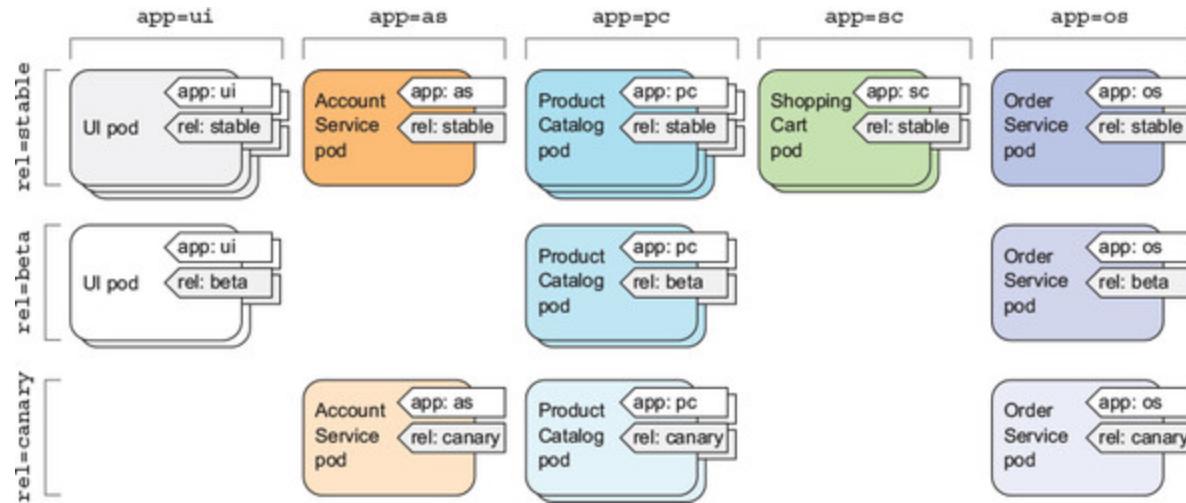


Picture from [Kubernetes in action](#)

Labels

Labels are key/value pair tags (like tags in AWS)

Can be used to query and organize resources



Picture from [Kubernetes in action](#)

Lab 4: Applying labels

Labels in the object definition

```
apiVersion: v1
kind: Pod
metadata:
  name: echo-server
  labels:
    env: training
    type: single_pod
spec:
  containers:
  - name: echo
    image: k8s.gcr.io/echoserver:1.4
```

Checking the labels

```
kubectl apply -f labs/03_labels.yaml
kubectl get pods --show-labels
```

File: labs/03_labels.yaml

Let's kubectl

Lab 4: Labeling our Pods

<https://github.com/aws-els-lin/eks/tree/master/labs/04-labels>

Links

Pods

- <https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>
- <https://kubernetes.io/docs/concepts/workloads/pods/pod/>

Kubernetes in Action

- <https://www.safaribooksonline.com/library/view/kubernetes-in-action/9781617293726/>

Demos

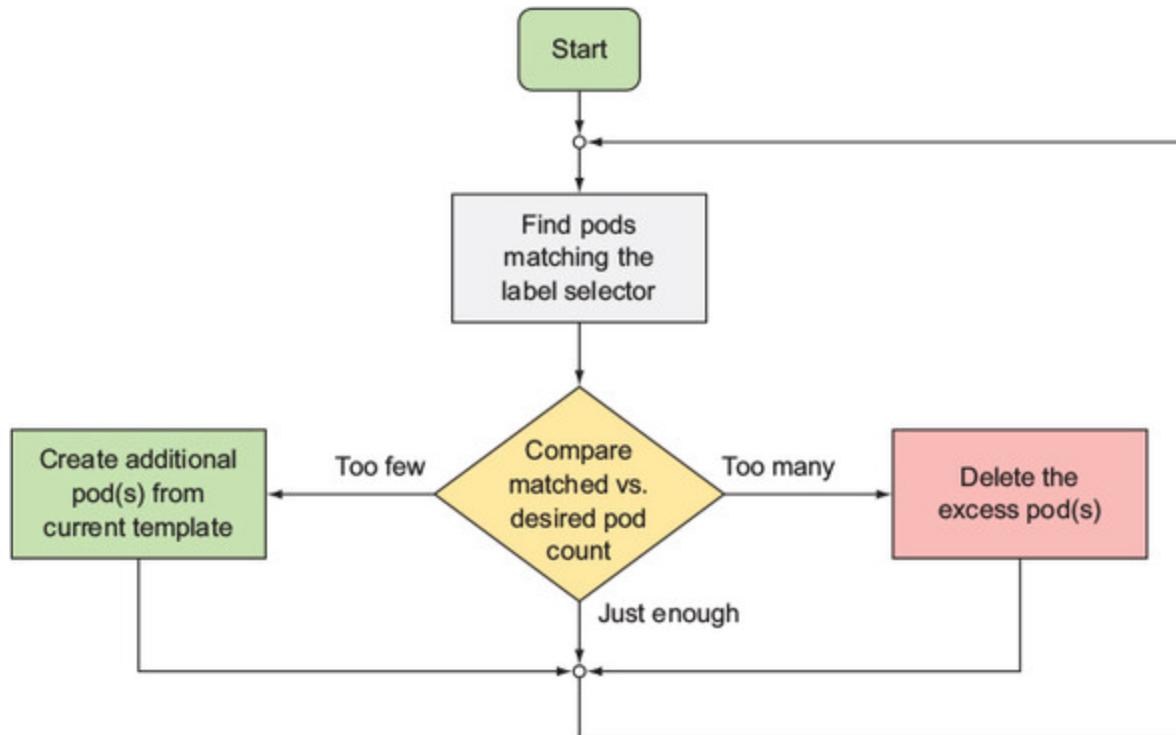
- <https://eksworkshop.com/>
- <https://github.com/kubernetes/contrib/blob/master/micro-demos/>
- <https://www.katacoda.com/>

Controllers

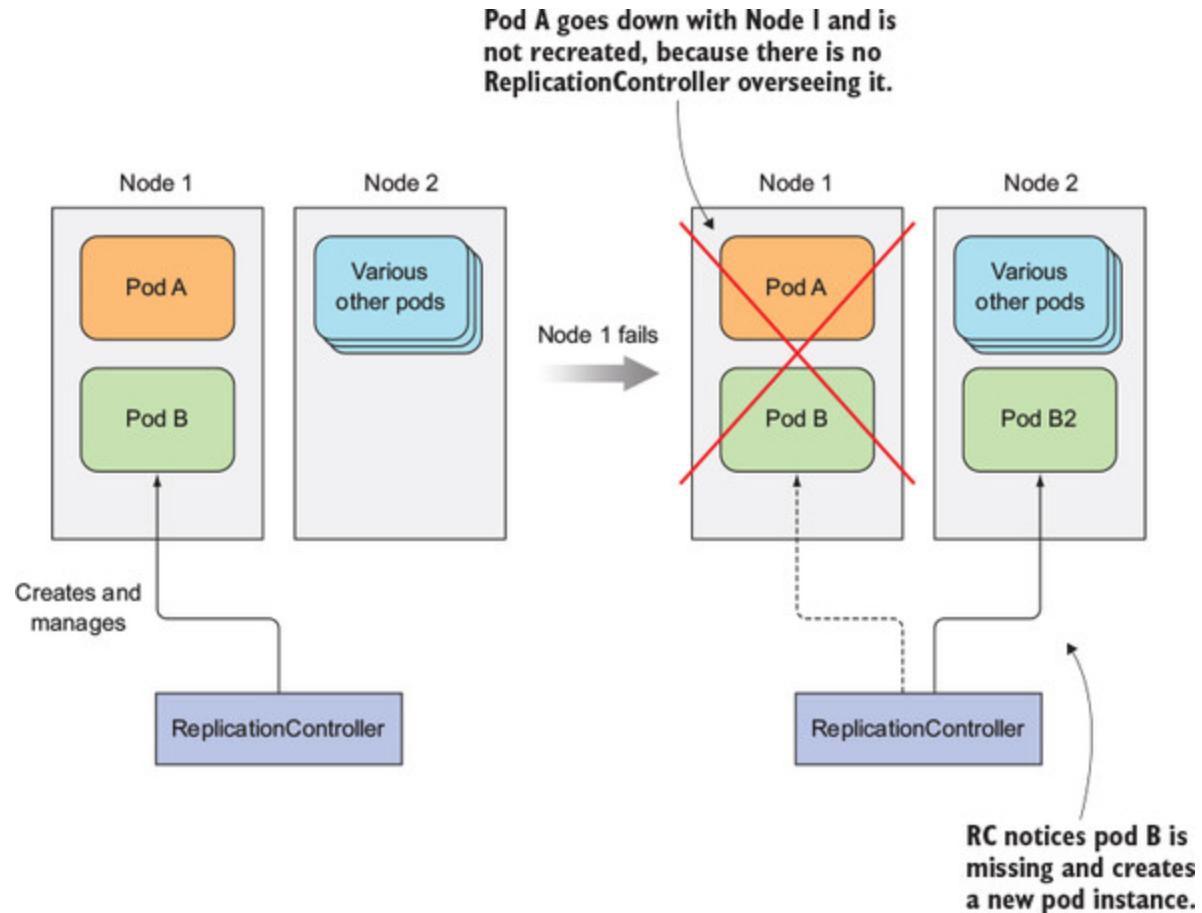
What controllers are and what they do?

- Resource responsible for managing pods
- Ensure pods are always running
- Replace missing and unhealthy pods
- Delete 'extra' pods
- Provide an easy way to scale the application
- Rely on Labels to account for the pods

What controllers are and what they do?



What controllers are and what they do?



Controllers

- Most commonly used controllers in Kubernetes:
 - **ReplicationController**
 - **ReplicaSet** the next generation of Replication Controllers
 - **Deployments** - preferred way to manage Replica Sets
 - **DaemonSet**
 - **Jobs**
 - **CronJobs**
 - **StatefulSets**

Deployments

Deployments

What is it?

- A Deployment controller provides declarative updates at controlled rate for Pods
- An easy way to deploy updates for existing applications
- Allows you to pause/resume deployments

Deployments

Comparing: Pods vs Deployment:

Single Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: web-server
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.7.9
```

Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-server-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
```

File: labs/04_deployment.yaml

Deployments

Lab 5: Create a Deployment

- Creating a Deployment

```
kubectl apply -f labs/04_deployment.yaml
```

- Checking the results

```
kubectl get deployments  
kubectl get pod  
kubectl get pod -l app=nginx
```

- Scale out the Deployment

```
kubectl scale deployment web-server-deployment --replicas=5
```

- Check the nginx server version

```
kubectl port-forward web-server-deployment-XXXXXX-YYYYYY 8080:80
```

Deployments

Lab 5: Scale the Deployment

- Scale in the deployment. Let's be frugal

```
kubectl edit deployment web-server-deployment  
# set spec.replicas to 3
```

- Checking the results

```
kubectl get pod  
# OR  
kubectl get pod -l app=nginx -L app
```

Deployments

Lab 5: Update the Deployment

- Update the deployment

```
kubectl set image deployment web-server-deployment nginx=httpd  
# OR  
kubectl edit deployment web-server-deployment  
# change spec.template.spec.containers.image to httpd
```

- Checking the results

```
kubectl rollout status deployment web-server-deployment  
# OR  
kubectl get pod  
# OR  
kubectl get pod -l app=nginx -L app
```

- Check the Nginx server version now

```
kubectl port-forward web-server-deployment-XXXXXX-YYYYY 8080:80
```

Deployments

Let's kubectl

Lab 5: Deployment

<https://github.com/aws-els-lin/tree/master/labs/05-Deployments>

Services

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

But why?

- Pods are ephemeral
- Pods' IPs are dynamic
- A single application might contain several Pods

Services

What's a service?

- Service is another layer on top of the pods
- Instead of connecting to the pods directly we connect to the service instead
- Very similar to a load balancer

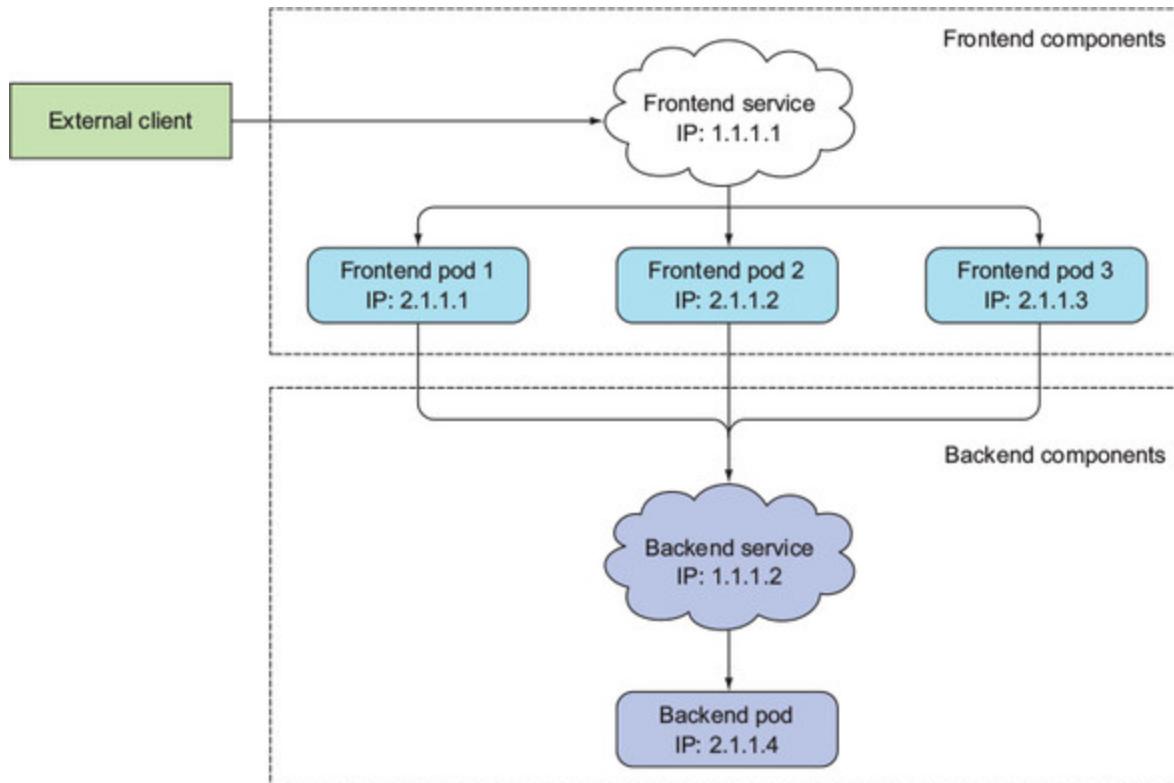
But why?

- Pods are ephemeral
- Pods' IPs are dynamic
- A single application might contain several Pods

So, how to reach an application?

Services

A **Service** is an abstraction layer which enables external traffic exposure, load balancing and service discovery



* from <https://kubernetes.io/>

Lab 6: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: web-app
spec:
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: web-server
```

File: labs/06-services

Lab 6: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: web-app
spec:
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: web-server
```

- Creating a service

```
kubectl apply -f service.yaml
```

File: labs/06-services

Lab 6: Create a Service

```
apiVersion: v1
kind: Service
metadata:
  name: web-app
spec:
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: web-server
```

- Creating a service

```
kubectl apply -f service.yaml
```

- Checking the results

```
kubectl get services
kubectl describe svc web-app
```

File: labs/06-services

Services

- There are different types of services:
 - ClusterIP is the default. Used for intra-cluster communication
 - LoadBalancer provisions a Load Balancer for you.

Connecting to your service:

```
kubectl edit svc web-app  
# change spec.type from 'ClusterIP' to 'LoadBalancer'
```

- Checking the results

```
kubectl get svc web-app
```

- Now you should get the LB URL from the EC2 console and open it in your browser.

Let's kubectl

Lab 6: Services

<https://github.com/aws-els-lin/eks/tree/master/labs/06-services>

Follow the steps in the `labs/06-services/README.md` file.

End of Day 1

Questions?

Thank you!